

# Classification II: evaluation & tuning

# Session learning objectives

By the end of the session, learners will be able to do the following:

- Describe what training, validation, and test data sets are and how they are used in classification.
- Split data into training, validation, and test data sets.
- Describe what a random seed is and its importance in reproducible data analysis.
- Set the random seed in Python using the `numpy.random.seed` function.
- Describe and interpret accuracy, precision, recall, and confusion matrices.

# Session learning objectives cont'd

By the end of the session, learners will be able to do the following:

- Evaluate classification accuracy, precision, and recall in Python using a test set, a single validation set, and cross-validation.
- Produce a confusion matrix in Python.
- Choose the number of neighbors in a K-nearest neighbors classifier by maximizing estimated cross-validation accuracy.
- Describe underfitting and overfitting, and relate it to the number of neighbors in K-nearest neighbors classification.
- Describe the advantages and disadvantages of the K-nearest neighbors classification algorithm.

# Evaluating performance

- Sometimes our classifier might make the wrong prediction.
- A classifier does not need to be right 100% of the time to be useful, though we don't want the classifier to make too many wrong predictions.
- How do we measure how “good” our classifier is?

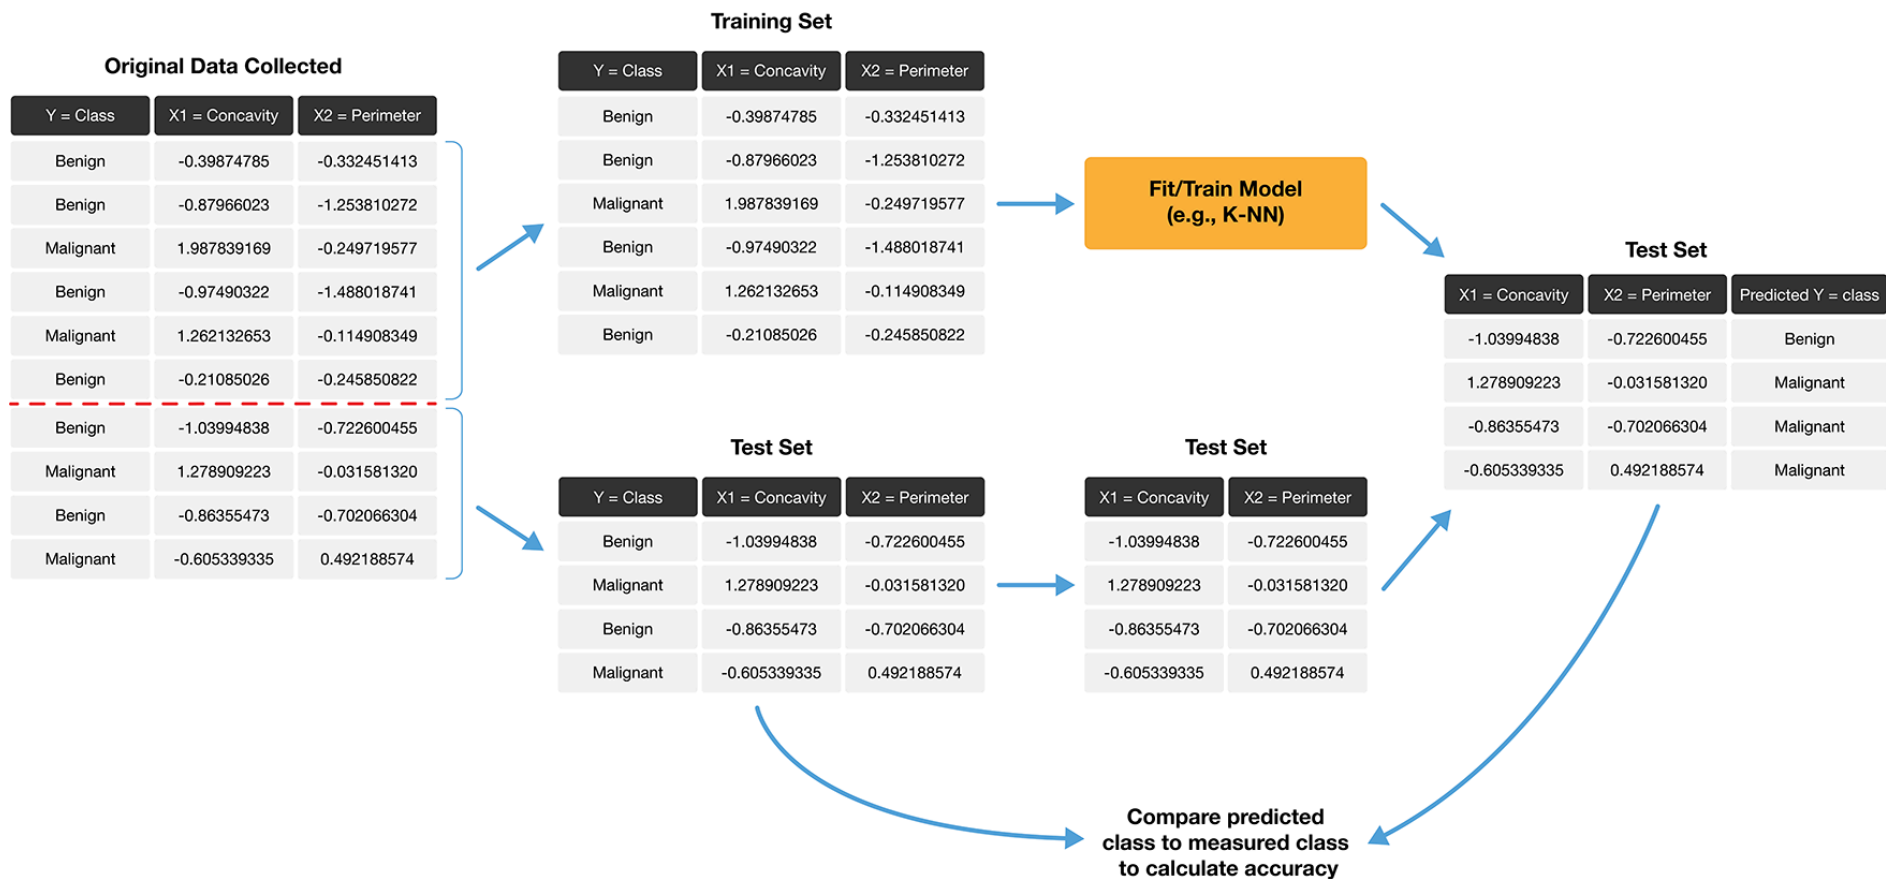
# Data splitting

- The trick is to split the data into a **training set** and **test set**.
- Only the **training set** when building the classifier.
- To evaluate the performance of the classifier, we first set aside the labels from the **test set**, and then use the classifier to predict the labels in the **test set**.
- If our predictions match the actual labels for the observations in the **test set**, then we have some confidence that our classifier might also accurately predict the class labels for new observations without known class labels.

# Splitting the data into training and testing sets

# Prediction accuracy

$$\text{accuracy} = \frac{\text{number of correct predictions}}{\text{total number of predictions}}$$



$$\text{accuracy} = \frac{\# \text{ correct predictions}}{\# \text{ total predictions}}$$

# Is knowing accuracy enough?

- Example accuracy calculation:

$$\text{accuracy} = \frac{\text{number of correct predictions}}{\text{total number of predictions}} = \frac{58}{65} = 0.892$$

- Prediction accuracy only tells us how often the classifier makes mistakes in general, but does not tell us anything about the *kinds* of mistakes the classifier makes.
- The **confusion matrix** tells a more complete story.



# Example confusion matrix for the breast cancer data:

	Predicted Malignant	Predicted Benign
Actually Malignant	1	3
Actually Benign	4	57

- **True Positive:** A malignant observation that was classified as malignant (top left).
- **False Positive:** A benign observation that was classified as malignant (bottom left).
- **True Negative:** A benign observation that was classified as benign (bottom right).
- **False Negative:** A malignant observation that was classified as benign (top right).

# Precision & recall

- *Precision* quantifies how many of the positive predictions the classifier made were actually positive.

$$\text{precision} = \frac{\text{number of correct positive predictions}}{\text{total number of positive predictions}}$$

- *Recall* quantifies how many of the positive observations in the test set were identified as positive.

$$\text{recall} = \frac{\text{number of correct positive predictions}}{\text{total number of positive test set observations}}$$

# Precision and recall for the breast cancer data set example

	Predicted Malignant	Predicted Benign
Actually Malignant	1	3
Actually Benign	4	57

$$\text{precision} = \frac{1}{1 + 4} = 0.20, \quad \text{recall} = \frac{1}{1 + 3} = 0.25$$

So even with an accuracy of 89%, the precision and recall of the classifier were both relatively low. For this data analysis context, recall is particularly important: if someone has a malignant tumor, we certainly want to identify it. A recall of just 25% would likely be unacceptable!

# Randomness and seeds

- Our data analyses will often involve the use of randomness
- We use randomness any time we need to make a decision in our analysis that needs to be fair, unbiased, and not influenced by human input (e.g., splitting into training and test sets).
- However, the use of randomness runs counter to one of the main tenets of good data analysis practice: reproducibility...
- The trick is that in Python—and other programming languages—randomness is not actually random! Instead, Python uses a random number generator that produces a sequence of numbers that are completely determined by a seed value.
- Once you set the seed value, everything after that point may look random, but is actually totally reproducible.

# Setting the seed in Python

Let's say we want to make a series object containing the integers from 0 to 9. And then we want to randomly pick 10 numbers from that list, but we want it to be reproducible.

```
1 import numpy as np
2 import pandas as pd
3
4 np.random.seed(1)
5
6 nums_0_to_9 = pd.Series([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
7
8 random_numbers1 = nums_0_to_9.sample(n=10).to_list()
9 random_numbers1
```

```
[2, 9, 6, 4, 0, 3, 1, 7, 8, 5]
```

You can see that `random_numbers1` is a list of 10 numbers from 0 to 9 that, from all appearances, looks random. If we run the `sample` method again, we will get a fresh batch of 10 numbers that also look random.

```
1 random_numbers2 = nums_0_to_9.sample(n=10).to_list()
2 random_numbers2
```

```
[9, 5, 3, 0, 8, 4, 2, 1, 6, 7]
```

# Setting the seed in Python (cont'd)

If we choose a different value for the seed—say, 4235—we obtain a different sequence of random numbers:

```
1 # seed was 1 on the last slide when we generated this list
2 random_numbers1
```

```
[2, 9, 6, 4, 0, 3, 1, 7, 8, 5]
```

```
1 np.random.seed(4235)
2 random_numbers1_different = nums_0_to_9.sample(n=10).to_list()
3
4 random_numbers1
5 random_numbers1_different
```


```
[6, 7, 2, 3, 5, 9, 1, 4, 0, 8]
```

# Back to the breast cancer data set example

```
1 # load packages
2 import altair as alt
3 import pandas as pd
4 from sklearn import set_config
5
6 # Output dataframes instead of arrays
7 set_config(transform_output="pandas")
8
9 # set the seed
10 np.random.seed(3)
11
12 # load data
13 cancer = pd.read_csv("data/wdbc_unscaled.csv")
14 # re-label Class "M" as "Malignant", and Class "B" as "Benign"
15 cancer["Class"] = cancer["Class"].replace({
16     "M" : "Malignant",
17     "B" : "Benign"
18 })
```

# Breast cancer data

```
1  perim_concav = alt.Chart(cancer).mark_circle().enc
2      x=alt.X("Smoothness").scale(zero=False),
3      y="Concavity",
4      color=alt.Color("Class").title("Diagnosis")
5  )
```





# Create the train / test split

- **Before** fitting any models, or doing exploratory data analysis, it is critical that you split the data into training and test sets.
- Typically, the training set is between 50% and 95% of the data, while the test set is the remaining 5% to 50%.
- The `train_test_split` function from `scikit-learn` handles the procedure of splitting the data for us.
- Use `shuffle=True` to remove the influence of order in the data set.
- Set the `stratify` parameter to be the response variable to ensure the same proportion of each class ends up in both the training and testing sets.

# Splitting the breast cancer data set

- Split the data so 75% are in the training set, and 25% in the test set
- Data are shuffled
- Split is stratified on the `Class` variable

```
1 from sklearn.model_selection import train_test_split
2
3 cancer_train, cancer_test = train_test_split(
4     cancer, train_size=0.75, stratify=cancer["Class"]
5 )
```

# Checking the splits

- We can use `.info()` to look at the splits
- Let's look at the training split (in practice you look at both)

```
1 cancer_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 426 entries, 196 to 296
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	ID	426 non-null	int64
1	Class	426 non-null	object
2	Radius	426 non-null	float64
3	Texture	426 non-null	float64
4	Perimeter	426 non-null	float64
5	Area	426 non-null	float64
6	Smoothness	426 non-null	float64
7	Compactness	426 non-null	float64
8	Concavity	426 non-null	float64
9	Concave_Points	426 non-null	float64
10	Symmetry	426 non-null	float64
11	Fractal_Dimension	426 non-null	float64

```
dtypes: float64(10), int64(1), object(1)
```

```
memory usage: 43.3+ KB
```

# Checking the splits

- We can use the `value_counts` method with the `normalize` argument set to `True` to find the percentage of malignant and benign classes in `cancer_train`.
- We can see our class proportions were roughly preserved when we split the data.

```
1 cancer_train["Class"].value_counts(normalize=True)
```

```
Class
Benign      0.626761
Malignant   0.373239
Name: proportion, dtype: float64
```

# Preprocessing with data splitting

- Many machine learning models are sensitive to the scale of the predictors, and even if not, comparison of importance of features for prediction after fitting requires scaling.
- When preprocessing the data (scaling is part of this), it is critical that we use **only the training set** in creating the mathematical function to do this.
- If this is not done, we will get overly optimistic test accuracy, as our test data will have influenced our model.
- After creating the preprocessing function, we can then apply it separately to both the training and test data sets.

# Preprocessing with `scikit-learn`

- `scikit-learn` helps us handle this properly as long as we wrap our analysis steps in a `Pipeline`.
- Specifically, we construct and prepare the preprocessor using `make_column_transformer`:

```
1 from sklearn.preprocessing import StandardScaler
2 from sklearn.compose import make_column_transformer, make_column_selector
3
4 cancer_preprocessor = make_column_transformer(
5     (StandardScaler(), ["Smoothness", "Concavity"]),
6 )
```

# Train the classifier

# Predict the labels in the test set

Now that we have a K-nearest neighbors classifier object, we can use it to predict the class labels for our test set:

```
1 cancer_test["predicted"] = knn_pipeline.predict(cancer_test[["Smoothness", "Concavity"]])
2 print(cancer_test[["ID", "Class", "predicted"]])
```

	ID	Class	predicted
116	864726	Benign	Malignant
146	869691	Malignant	Malignant
..	...	...	...
281	8912055	Benign	Benign
15	84799002	Malignant	Malignant

```
[143 rows x 3 columns]
```



# Evaluate performance

To evaluate the model, we will look at:

- accuracy
- precision
- recall
- confusion matrix
- compare to baseline model (majority classifier)

All of these together, will help us develop a fuller picture of how the model is performing, as opposed to only evaluating the model based on a single metric or table.

# Accuracy, precision and recall

```
1 knn_pipeline.score(  
2     cancer_test[["Smoothness", "Concavity"]],  
3     cancer_test["Class"]  
4 )
```

0.8951048951048951

```
1 from sklearn.metrics import recall_score, precision_score  
2  
3 precision_score(  
4     y_true=cancer_test["Class"],  
5     y_pred=cancer_test["predicted"],  
6     pos_label="Malignant"  
7 )
```

np.float64(0.8275862068965517)

```
1 recall_score(  
2     y_true=cancer_test["Class"],  
3     y_pred=cancer_test["predicted"],  
4     pos_label="Malignant"  
5 )
```

np.float64(0.9056603773584906)

# Confusion matrix

- We can look at the *confusion matrix* for the classifier using the `crosstab` function from `pandas`.
- The `crosstab` function takes two arguments: the actual labels first, then the predicted labels second.
- Note that `crosstab` orders its columns alphabetically, but the positive label is still `Malignant`, even if it is not in the top left corner as in the table shown earlier.

```
1 pd.crosstab(  
2     cancer_test["Class"],  
3     cancer_test["predicted"]  
4 )
```

predicted	Benign	Malignant
Class		
Benign	80	10
Malignant	5	48

# Critically analyze performance

- Is 90% accuracy, a precision of 83% and a recall of 91% good enough?
- To get a sense of scale, we often compare our model to a baseline model. In the case of classification, this would be the majority classifier (*always* guesses the majority class label from the training data).
- For the breast cancer training data, the baseline classifier's accuracy would be 63%

```
1 cancer_train["Class"].value_counts(normalize=True)
```

```
Class
Benign      0.626761
Malignant   0.373239
Name: proportion, dtype: float64
```

- So we do see that our model is doing a LOT better than the baseline, which is great, but considering our application domain is in cancer diagnosis, we still have a ways to go...
- Analyzing model performance really depends on your application!

# Tuning the classifier

- Most predictive models in statistics and machine learning have parameters (a number you have to pick in advance that determines some aspect of how the model behaves).
- For our working example,  $K$ -nearest neighbors classification algorithm,  $K$  is a parameter that we have to pick that determines how many neighbors participate in the class vote.
- How do we choose  $K$ , or any parameter for other models?
- **Data splitting!**

# Validation set

- Cannot use the test set to choose the parameter!
- But we can split the training set into two partitions, a training set and a validation set.
- For each parameter value we want to assess, we can fit on the training set, and evaluate on the validation set.
- Then after we find the best value for our parameter, we can refit the model with the best parameter on the entire training set and then evaluate our model on the test set.

# Can we do better?

- Depending on how we split the data into the training and validation sets, we might get a lucky split (or an unlucky one) that doesn't give us a good estimate of the model's true accuracy.
- In many cases, we can do better by making many splits, and averaging the accuracy scores to get a better estimate.
- We call this cross-validation.

# Cross-validation

An example of 5-fold cross-validation:



$$\text{cross validation accuracy} = \frac{\text{accuracy}_1 + \text{accuracy}_2 + \text{accuracy}_3 + \text{accuracy}_4 + \text{accuracy}_5}{\text{\#folds}}$$



# Cross-validation in `scikit-learn`

# Parameter value selection

- Since cross-validation helps us evaluate the accuracy of our classifier, we can use cross-validation to calculate an accuracy for each value of our parameter, here  $K$ , in a reasonable range,
- Then we pick the value of  $K$  that gives us the best accuracy, and refit the model with our parameter on the training data, and then evaluate on the test data.
- The `scikit-learn` package collection provides built-in functionality, named `GridSearchCV`, to automatically handle the details for us.

# Parameter value selection

```
1 knn = KNeighborsClassifier() #don't specify the number of neighbours
2 cancer_tune_pipe = make_pipeline(cancer_preprocessor, knn)
3
4 parameter_grid = {
5     "kneighborsclassifier__n_neighbors": range(1, 100, 5),
6 }
7
8 from sklearn.model_selection import GridSearchCV
9
10 cancer_tune_grid = GridSearchCV(
11     estimator=cancer_tune_pipe,
12     param_grid=parameter_grid,
13     cv=10
14 )
```

# Parameter value selection

Now we use the fit method on the `GridSearchCV` object to begin the tuning process.

```
1 cancer_tune_grid.fit(  
2     cancer_train[["Smoothness", "Concavity"]],  
3     cancer_train["Class"]  
4 )  
5 accuracies_grid = pd.DataFrame(cancer_tune_grid.cv  
6 accuracies_grid.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 20 entries, 0 to 19  
Data columns (total 19 columns):  
#   Column                                     Non-Null  
Count  Dtype  
---  -  
-----  
0    mean_fit_time                             20 non-  
null    float64  
1    std_fit_time                               20 non-  
null    float64  
2    mean_score_time                           20 non-  
null    float64  
3    std_score_time                             20 non-  
null    float64  
4    param_kneighborsclassifier__n_neighbors  20 non-  
null    int64  
5    params                                     20 non-  
null    object  
6    split0_test_score                         20 non-  
null    float64  
7    split1_test_score                         20 non-
```

# Parameter value selection


```
1 accuracies_grid["sem_test_score"] = accuracies_grid["std_test_score"] / 10**(1/2)
2 accuracies_grid = (
3     accuracies_grid[[
4         "param_kneighborsclassifier__n_neighbors",
5         "mean_test_score",
6         "sem_test_score"
7     ]]
8     .rename(columns={"param_kneighborsclassifier__n_neighbors": "n_neighbors"})
9 )
10 print(accuracies_grid)
```

	n_neighbors	mean_test_score	sem_test_score
0	1	0.845127	0.019966
1	6	0.873200	0.015680
..	...	...	...
18	91	0.875581	0.012967
19	96	0.875581	0.008193

[20 rows x 3 columns]

# Visualize paramter value selection

```
1 accuracy_vs_k = (  
2     alt.Chart(accuracies_grid)  
3     .mark_line(point=True)  
4     .encode(  
5         x=alt.X("n_neighbors")  
6         .title("Neighbors"),  
7         y=alt.Y("mean_test_score")  
8         .scale(zero=False)  
9         .title("Accuracy estimate")  
10    )  
11 )
```



# Best parameter value

We can also obtain the number of neighbours with the highest accuracy programmatically by accessing the `best_params_` attribute of the fit `GridSearchCV` object.


```
1 cancer_tune_grid.best_params_  
{'kneighborsclassifier__n_neighbors': 36}
```

# Best parameter value

Do we use  $K = 36$ ?

Generally, when selecting a parameters, we are looking for a value where:


- we get roughly optimal accuracy
- changing the value to a nearby one doesn't change the accuracy too much
- the cost of training the model is not prohibitive







# Under/Overfitting



- What happens if we keep increasing the number of neighbors  $K$ ?
- The cross-validation accuracy estimate actually starts to decrease!





**Diagnosis**

- Benign
- Malignant



# Evaluating on the test set

- Before we evaluate on the test set, we need to refit the model using the best parameter(s) on the entire training set
- Luckily, `scikit-learn` does it for us automatically!
- To make predictions and assess the estimated accuracy of the best model on the test data, we can use the `score` and `predict` methods of the fit `GridSearchCV` object.

```
1 cancer_test["predicted"] = cancer_tune_grid.predict(  
2     cancer_test[["Smoothness", "Concavity"]]  
3 )
```

# Evaluating on the test set

We can then pass those predictions to the `precision`, `recall`, and `crosstab` functions to assess the estimated precision and recall, and print a confusion matrix.

```
1 cancer_tune_grid.score(  
2     cancer_test[["Smoothness", "Concavity"]],  
3     cancer_test["Class"]  
4 )
```

```
0.9090909090909091
```

```
1 precision_score(  
2     y_true=cancer_test["Class"],  
3     y_pred=cancer_test["predicted"],  
4     pos_label='Malignant'  
5 )
```

```
np.float64(0.8846153846153846)
```

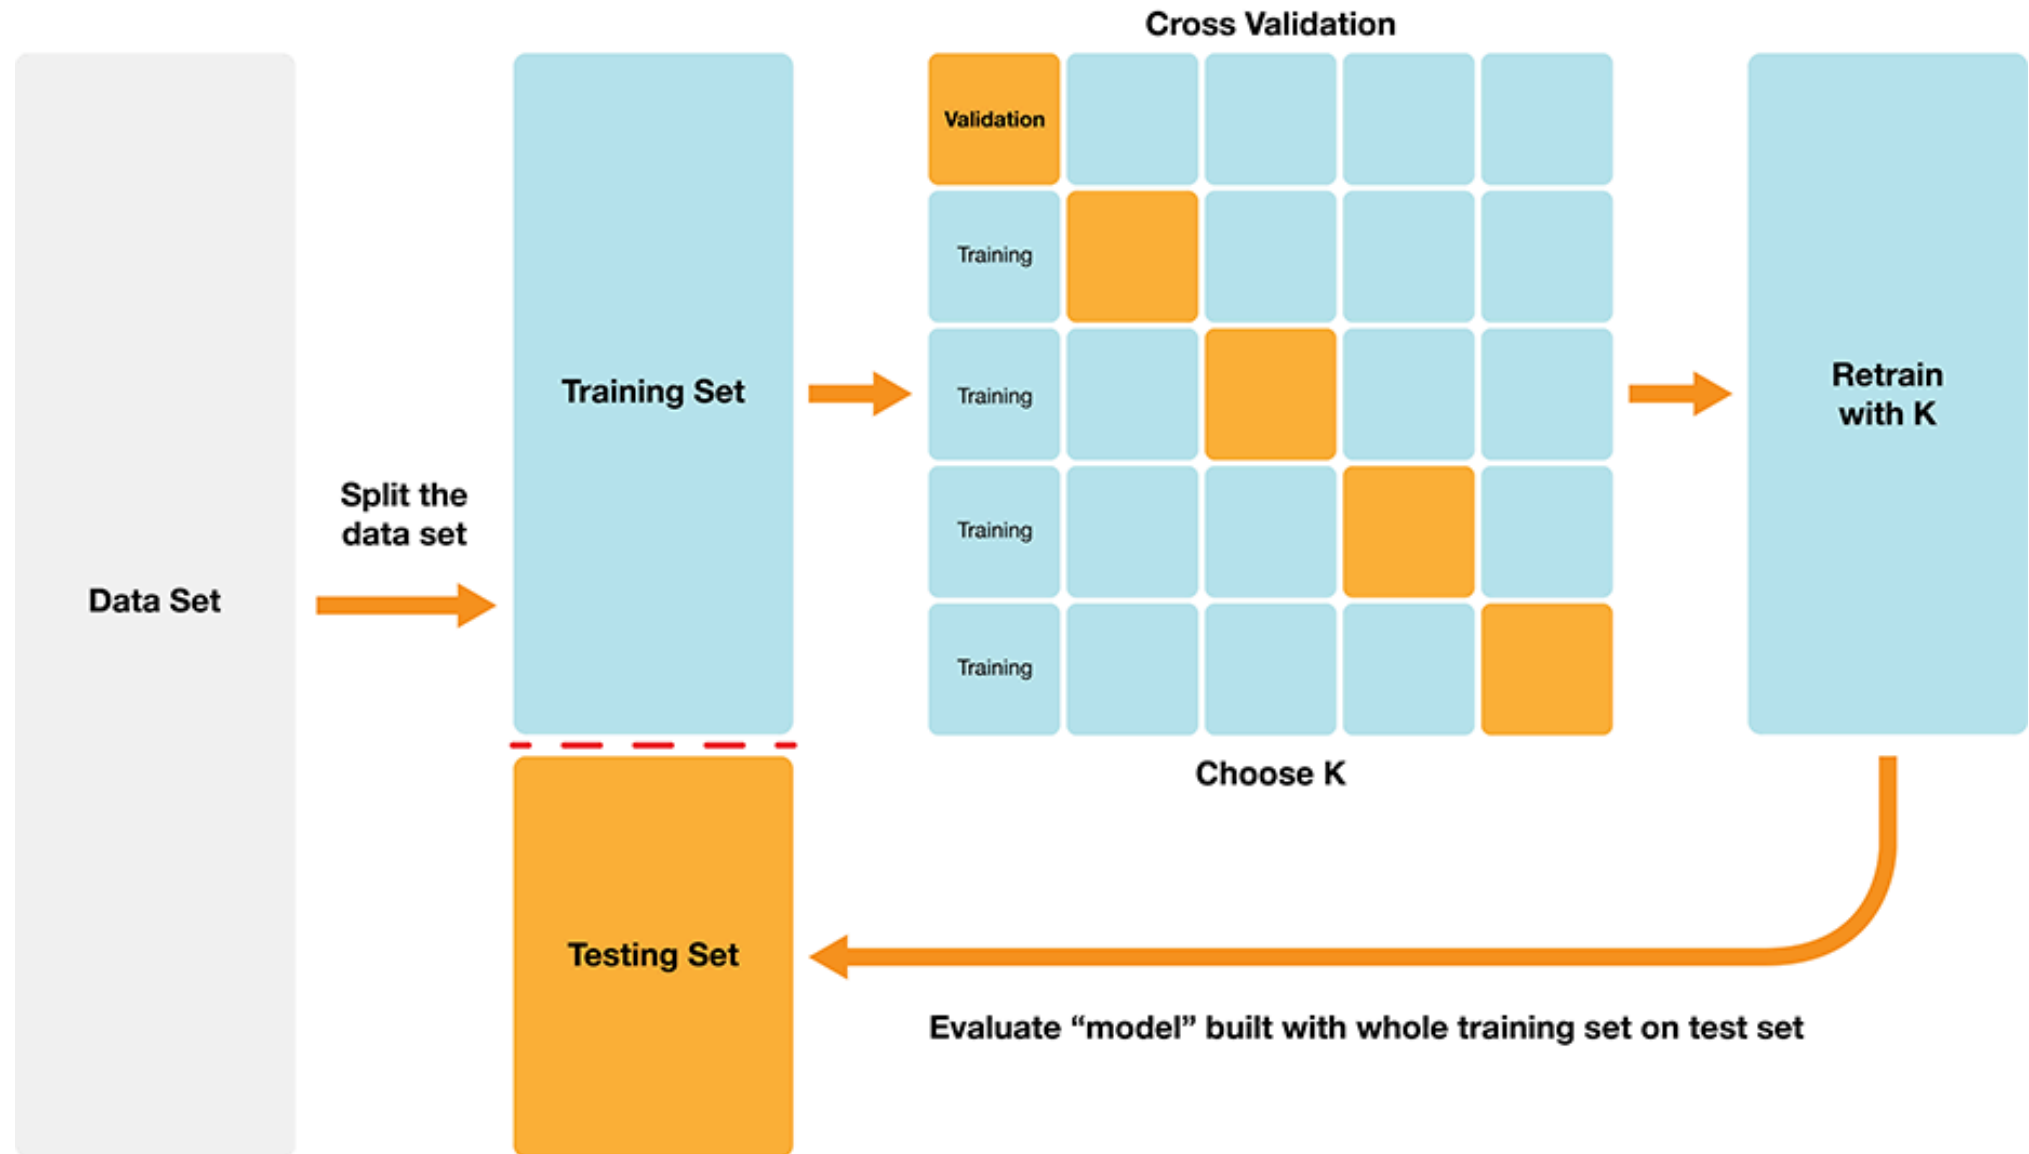
```
1 recall_score(  
2     y_true=cancer_test["Class"],  
3     y_pred=cancer_test["predicted"],  
4     pos_label='Malignant'  
5 )
```

```
np.float64(0.8679245283018868)
```

```
1 conf_matrix = pd.crosstab(  
2     cancer_test["Class"],  
3     cancer_test["predicted"]  
4 )  
5 print(conf_matrix)
```

predicted	Benign	Malignant
Class		
Benign	84	6
Malignant	7	46

# Summary

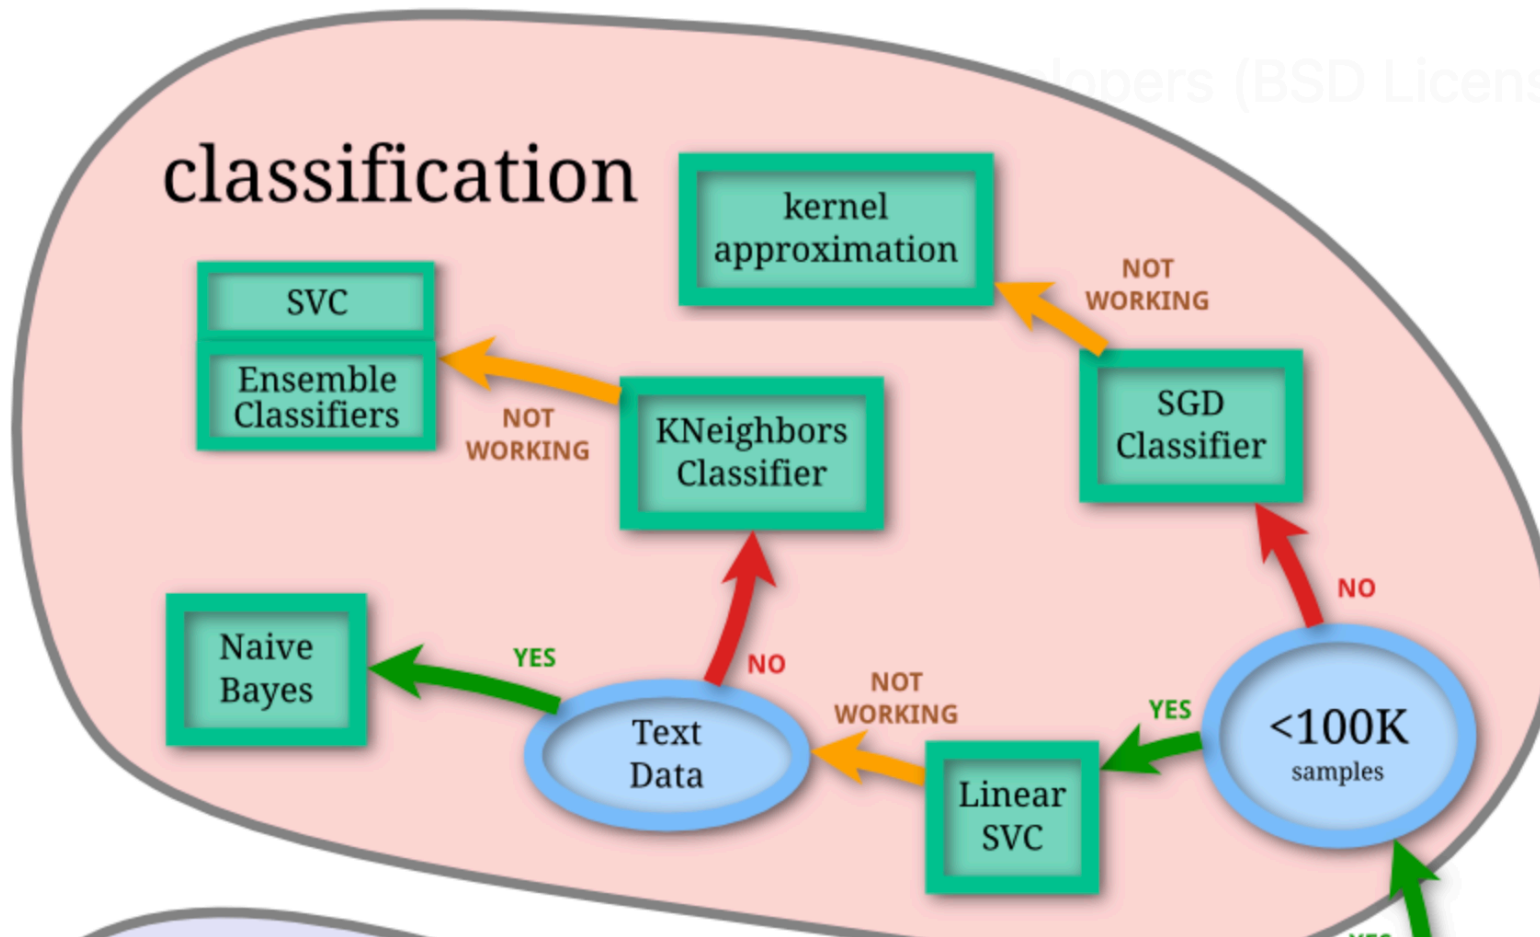


# K-nearest neighbors classification algorithm

**Strengths:** K-nearest neighbors classification ::: {.nonincremental} 1. is a simple, intuitive algorithm, 2. requires few assumptions about what the data must look like, and 3. works for binary (two-class) and multi-class (more than 2 classes) classification problems. :::

**Weaknesses:** K-nearest neighbors classification ::: {.nonincremental} 1. becomes very slow as the training data gets larger, 2. may not perform well with a large number of predictors, and 3. may not perform well when classes are imbalanced. :::

# Other classification algorithms



*scikit-learn* classification documentation: [https://scikit-learn.org/stable/supervised\\_learning.html](https://scikit-learn.org/stable/supervised_learning.html)

# Additional resources

- The [Classification II: evaluation & tuning](#) chapter of Data Science: A First Introduction (Python Edition) by Tiffany Timbers, Trevor Campbell, Melissa Lee, Joel Ostblom, Lindsey Heagy contains all the content presented here with a detailed narrative.
- The [scikit-learn website](#) is an excellent reference for more details on, and advanced usage of, the functions and packages in the past two chapters. Aside from that, it also offers many useful [tutorials](#) to get you started.
- *[An Introduction to Statistical Learning](#)* {cite:p}james2013introduction provides a great next step in the process of learning about classification. Chapter 4 discusses additional basic techniques for classification that we do not cover, such as logistic regression, linear discriminant analysis, and naive Bayes.



# References

Evelyn Martin Lansdowne Beale, Maurice George Kendall, and David Mann. The discarding of variables in multivariate analysis. *Biometrika*, 54(3-4):357–366, 1967.

Norman Draper and Harry Smith. *Applied Regression Analysis*. Wiley, 1966.

M. Eforymson. Stepwise regression—a backward and forward look. In *Eastern Regional Meetings of the Institute of Mathematical Statistics*. 1966.

Ronald Hocking and R. N. Leslie. Selection of the best subset in regression analysis. *Technometrics*, 9(4):531–540, 1967.

Gareth James, Daniela Witten, Trevor Hastie, and Robert Tibshirani. *An Introduction to Statistical Learning*. Springer, 1st edition, 2013. URL: <https://www.statlearning.com/>.

Wes McKinney. *Python for data analysis: Data wrangling with Pandas, NumPy, and IPython*. ” O’Reilly Media, Inc.”, 2012.

William Nick Street, William Wolberg, and Olvi Mangasarian. Nuclear feature extraction for breast tumor diagnosis. In *International Symposium on Electronic Imaging: Science and Technology*. 1993.