

Regression

K-NN and Linear Regression

Session learning objectives: KNN

Session learning objective: Linear Regression

- Use Python to fit simple and multivariable linear regression models on training data.
- Evaluate the linear regression model on test data.
- Compare and contrast predictions obtained from K-nearest neighbors regression to those obtained using linear regression from the same data set.
- Describe how linear regression is affected by outliers and multicollinearity.

The regression problem

- Predictive problem
- Use past information to predict future observations
- Predict *numerical* values instead of *categorical* values

Examples:

- Race time in the Boston marathon
- size of a house to predict its sale price

Regression Methods

In this workshop:

- K-nearest neighbors
- Linear regression

Classification similarities to regression

Concepts from classification map over to the setting of regression

- Predict a new observation's response variable based on past observations
- Split the data into training and test sets
- Use cross-validation to evaluate different choices of model parameters

Difference

Predicting numerical variables instead of categorical variables

Explore a data set

932 real estate transactions in Sacramento, California

Can we use the size of a house in the Sacramento, CA area to predict its sale price?

Data and package setup

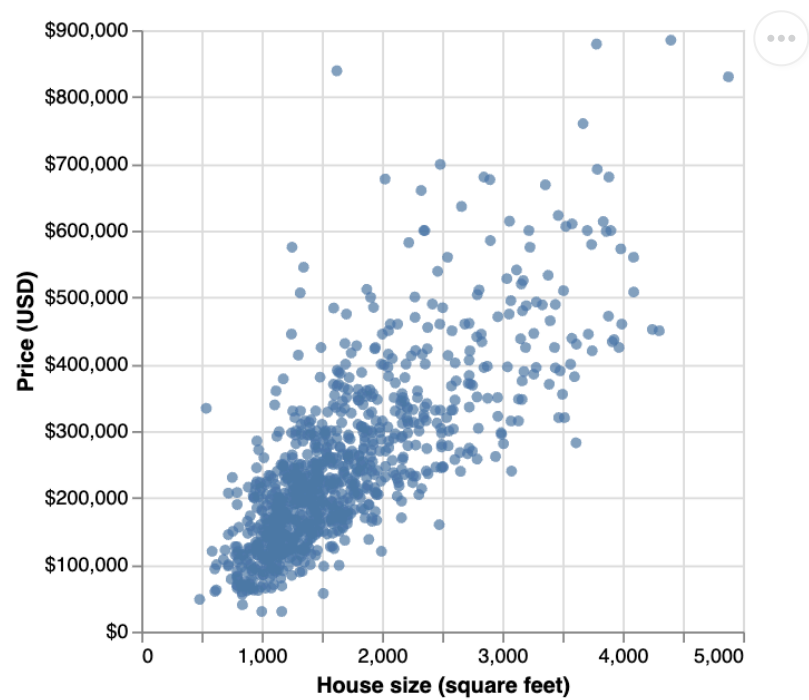
```
1 import altair as alt
2 import numpy as np
3 import pandas as pd
4 import plotly.express as px
5 import plotly.graph_objects as go
6 from sklearn.model_selection import GridSearchCV, train_test_split
7 from sklearn.compose import make_column_transformer
8 from sklearn.pipeline import make_pipeline
9 from sklearn.preprocessing import StandardScaler
10 from sklearn import set_config
11
12 # Output dataframes instead of arrays
13 set_config(transform_output='pandas')
14
15 sacramento = pd.read_csv('data/sacramento.csv')
16 print(sacramento)
```

	city	zip	beds	baths	sqft	type	price \
0	SACRAMENTO	z95838	2	1.0	836	Residential	59222
1	SACRAMENTO	z95823	3	1.0	1167	Residential	68212
..
930	ELK_GROVE	z95758	4	2.0	1685	Residential	235301
931	EL_DORADO_HILLS	z95762	3	2.0	1362	Residential	235738

	latitude	longitude
0	38.631913	-121.434879
1	38.478902	-121.431028
..
930	38.417000	-121.397424
931	38.655245	-121.075915

[932 rows x 9 columns]

Price vs Sq.Ft



K-nearest neighbors regression

```
1 # look at a small sample of data
2 np.random.seed(10)
3
4 small_sacramento = sacramento.sample(n=30)
5 print(small_sacramento)
```

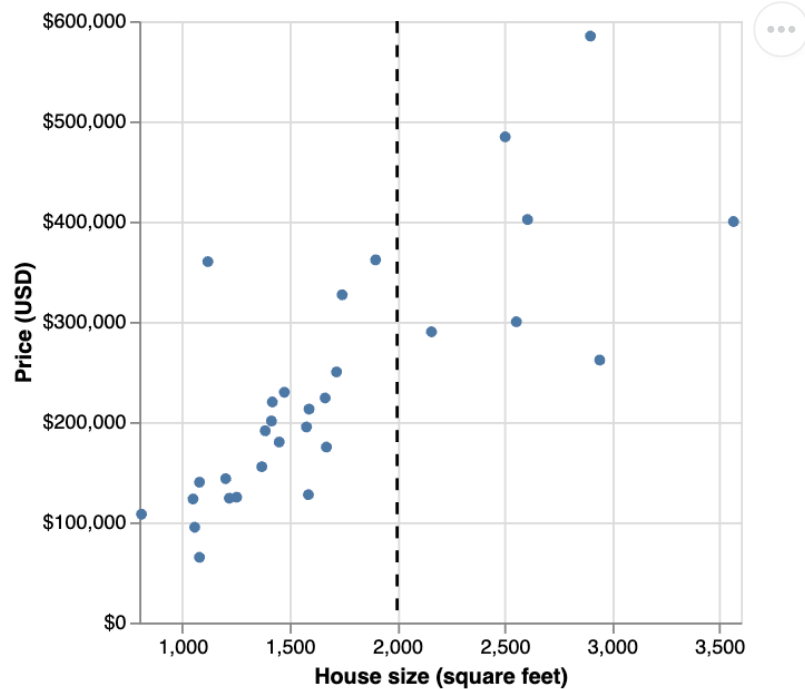
	city	zip	beds	baths	sqft	type	price	latitude \
538	ELK_GROVE	z95758	3	3.0	2503	Residential	484500	38.409689
304	ROCKLIN	z95765	4	2.0	2607	Residential	402000	38.805749
..
559	SACRAMENTO	z95817	2	1.0	1080	Residential	65000	38.544162
917	SACRAMENTO	z95834	3	2.0	1665	Residential	224000	38.631026

	longitude
538	-121.446059
304	-121.280931
..	...
559	-121.460652
917	-121.501879

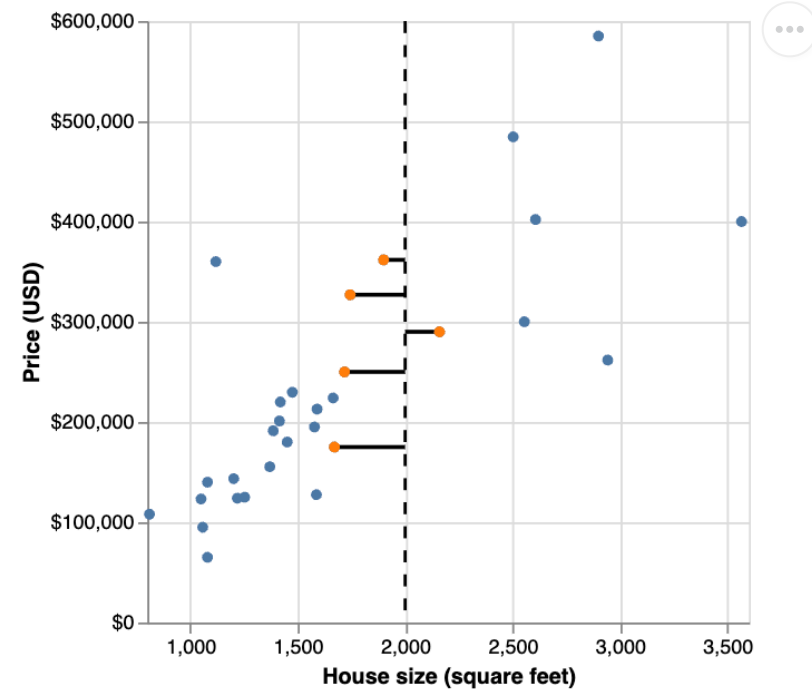
[30 rows x 9 columns]

Sample: Example

House price of 2000



5 closest neighbors



Sample: Predict

5 closest points

```
1 small_sacramento['dist'] = (2000 - small_sacramento['sqft']).abs()  
2 nearest_neighbors = small_sacramento.nsmallest(5, 'dist')  
3 print(nearest_neighbors)
```

	city	zip	beds	baths	sqft	type	price	\
298	SACRAMENTO	z95823	4	2.0	1900	Residential	361745	
718	ANTELOPE	z95843	4	2.0	2160	Residential	290000	
748	ROSEVILLE	z95678	3	2.0	1744	Residential	326951	
252	SACRAMENTO	z95835	3	2.5	1718	Residential	250000	
211	RANCHO_CORDOVA	z95670	3	2.0	1671	Residential	175000	

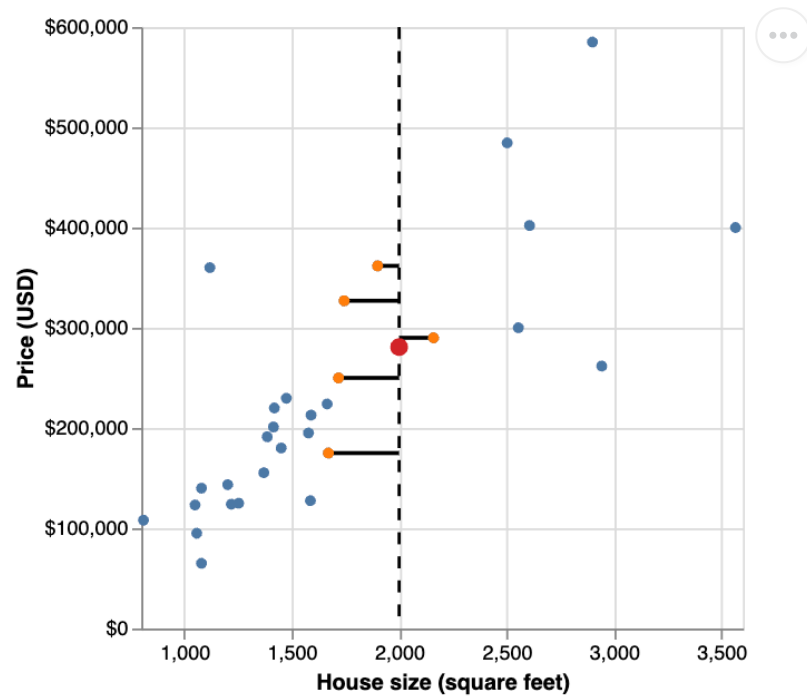
	latitude	longitude	dist
298	38.487409	-121.461413	100
718	38.704554	-121.354753	160
748	38.771917	-121.304439	256
252	38.676658	-121.528128	282
211	38.591477	-121.315340	329

Average of nearest points

```
1 prediction = nearest_neighbors['price'].mean()  
2 print(prediction)
```

280739.2

Sample: Visualize new prediction



```
1 print(prediction)
```

280739.2

Training, evaluating, and tuning the model

```
1 np.random.seed(1)
2
3 sacramento_train, sacramento_test = train_test_split(
4     sacramento, train_size=0.75
5 )
```

Note

We are not specifying the stratify argument. The `train_test_split()` function cannot stratify on a quantitative variable

Metric: RMS(P)E

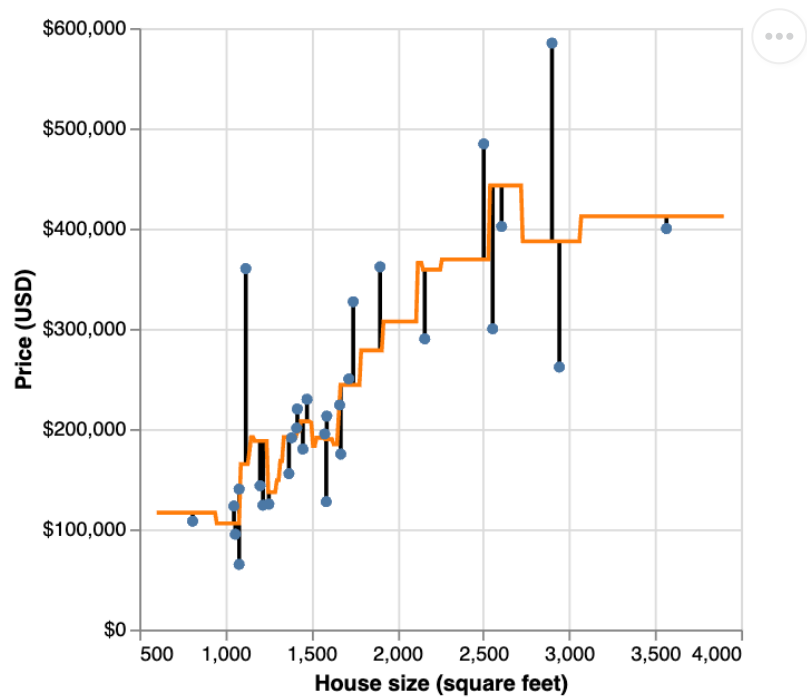
Root Mean Square (Prediction) Error

$$\text{RMSPE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

where:

- n is the number of observations,
- y_i is the observed value for the i^{th} observation, and
- \hat{y}_i is the forecasted/predicted value for the i^{th} observation.

Metric: Visualize



RMSPE vs RMSE

Root Mean Square (Prediction) Error

- RMSPE: the error calculated on the non-training dataset
- RMSE: the error calculated on the training dataset

This notation is a statistics distinction, you will most likely see RMSPE written as RMSE.

Pick the best k: `GridSearchCV()`

We'll use cross validation to try out many different values of `k`

```
1 # import the K-NN regression model
2 from sklearn.neighbors import KNeighborsRegressor
3
4 # preprocess the data, make the pipeline
5 sacr_preprocessor = make_column_transformer((StandardScaler(), ['sqft']))
6 sacr_pipeline = make_pipeline(sacr_preprocessor, KNeighborsRegressor())
7
8 # create the 5-fold GridSearchCV object
9 param_grid = {
10     'kneighborsregressor__n_neighbors': range(1, 201, 3),
11 }
12 sacr_gridsearch = GridSearchCV(
13     estimator=sacr_pipeline,
14     param_grid=param_grid,
15     cv=5,
16     scoring='neg_root_mean_squared_error', # we will deal with this later
17 )
```

Pick the best k: fit the CV models

```
1 # fit the GridSearchCV object
2 sac_gridsearch.fit(
3     sacramento_train[['sqft']], # A single-column data frame
4     sacramento_train['price'], # A series
5 )
6
7 # Retrieve the CV scores
8 sacr_results = pd.DataFrame(sacr_gridsearch.cv_results_)
9 sacr_results['sem_test_score'] = sacr_results['std_test_score'] / 5 ** (1 / 2)
10 sacr_results = sacr_results[
11     [
12         'param_kneighborsregressor__n_neighbors',
13         'mean_test_score',
14         'sem_test_score',
15     ]
16 ].rename(columns={'param_kneighborsregressor__n_neighbors': 'n_neighbors'})
17 print(sacr_results)
```

	n_neighbors	mean_test_score	sem_test_score
0	1	-117365.988307	2715.383001
1	4	-93956.523683	2466.200227
..
65	196	-93671.588088	2473.312705
66	199	-93986.752272	2473.048651

[67 rows x 3 columns]

Look at the CV Results

```
1 print(sacr_results)
```

	n_neighbors	mean_test_score	sem_test_score
0	1	-117365.988307	2715.383001
1	4	-93956.523683	2466.200227
..
65	196	-93671.588088	2473.312705
66	199	-93986.752272	2473.048651

[67 rows x 3 columns]

- **n_neighbors**: values of K
- **mean_test_score**: RMSPE estimated via cross-validation (it's negative!)
- **sem_test_score**: standard error of our cross-validation RMSPE estimate (how uncertain we are in the mean value)

```
1 sacr_results['mean_test_score'] = -sacr_results['mean_test_score']  
2 print(sacr_results)
```

	n_neighbors	mean_test_score	sem_test_score
0	1	117365.988307	2715.383001
1	4	93956.523683	2466.200227
..
65	196	93671.588088	2473.312705
66	199	93986.752272	2473.048651

[67 rows x 3 columns]

Best k

take the *minimum* RMSPE to find the best setting for the number of neighbors

```
1 best_k_sacr = sacr_results["n_neighbors"][sacr_results["mean_test_score"].idxmin()]
2 best_cv_RMSPE = min(sacr_results["mean_test_score"])
```

Best k:

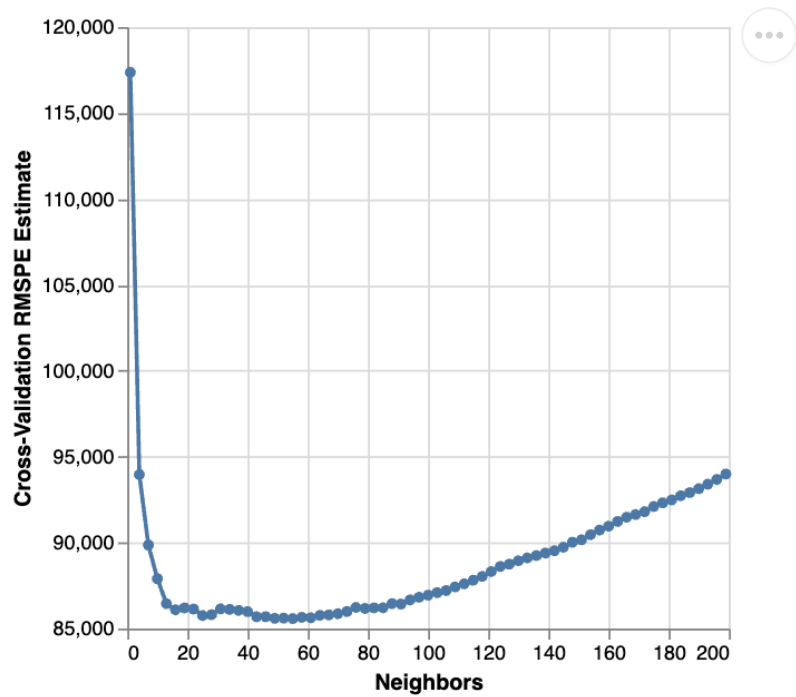
```
1 best_k_sacr
```

`np.int64(55)`

```
1 best_cv_RMSPE
```

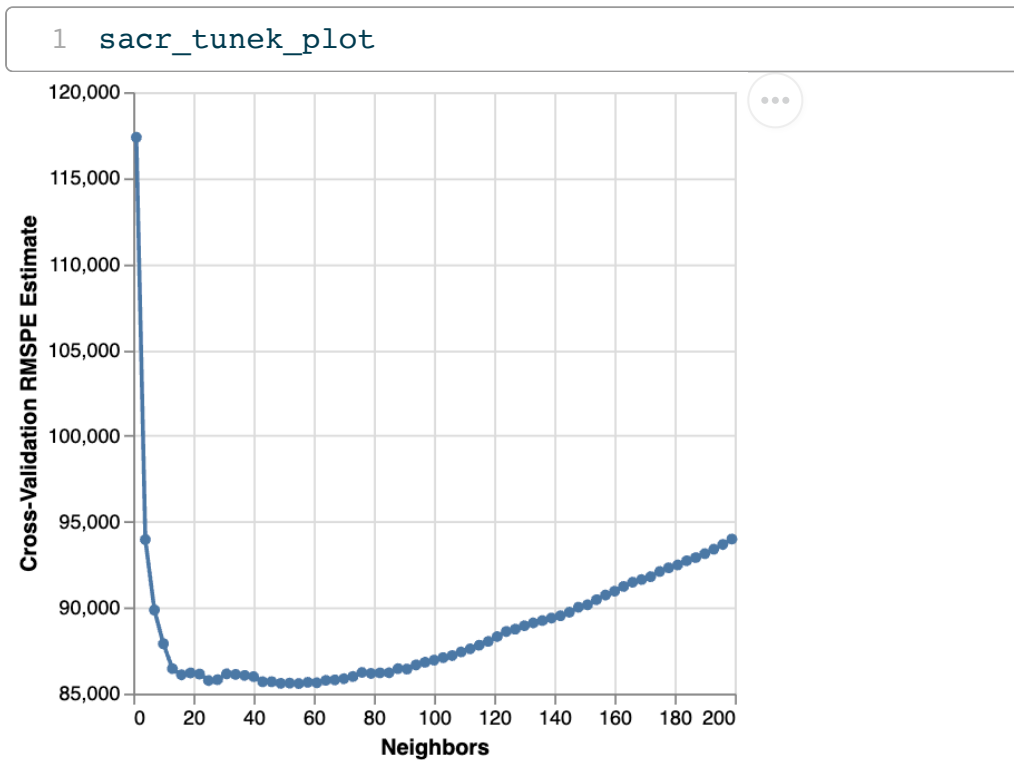
`85578.21347207513`

Best k: Visualize



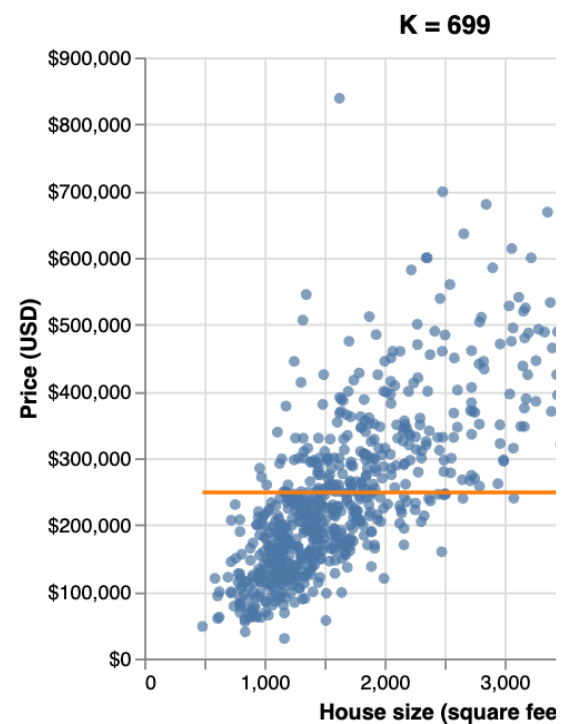
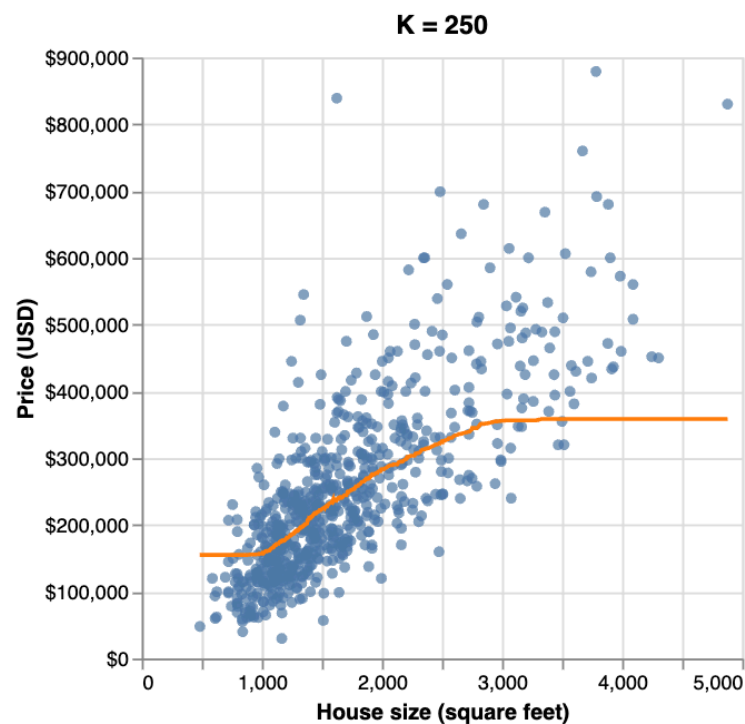
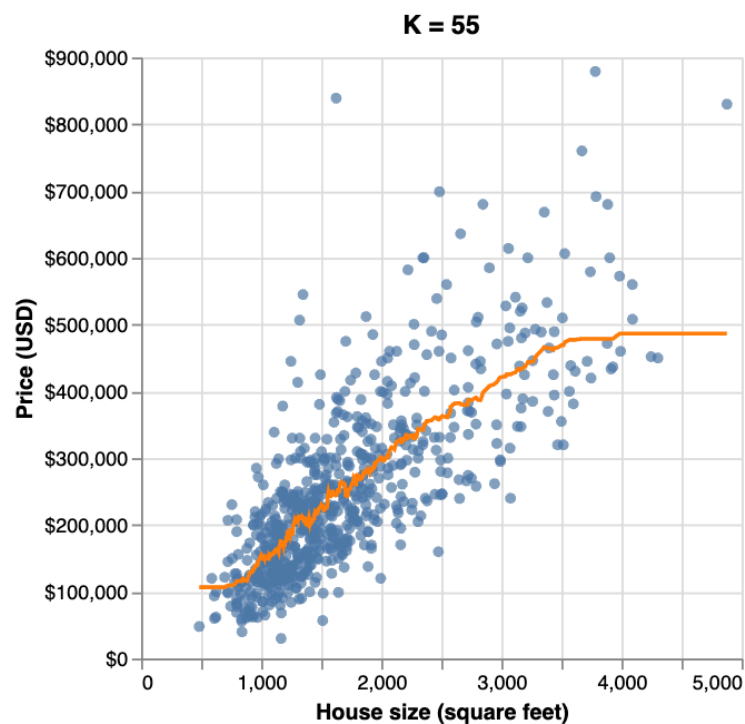
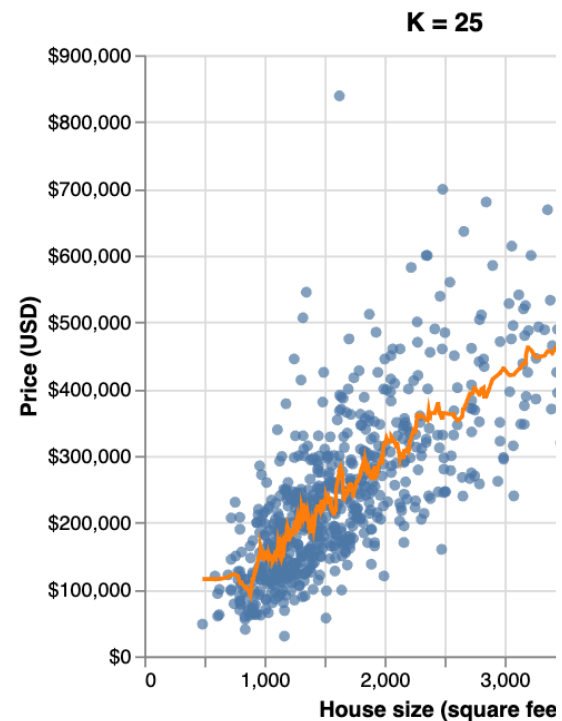
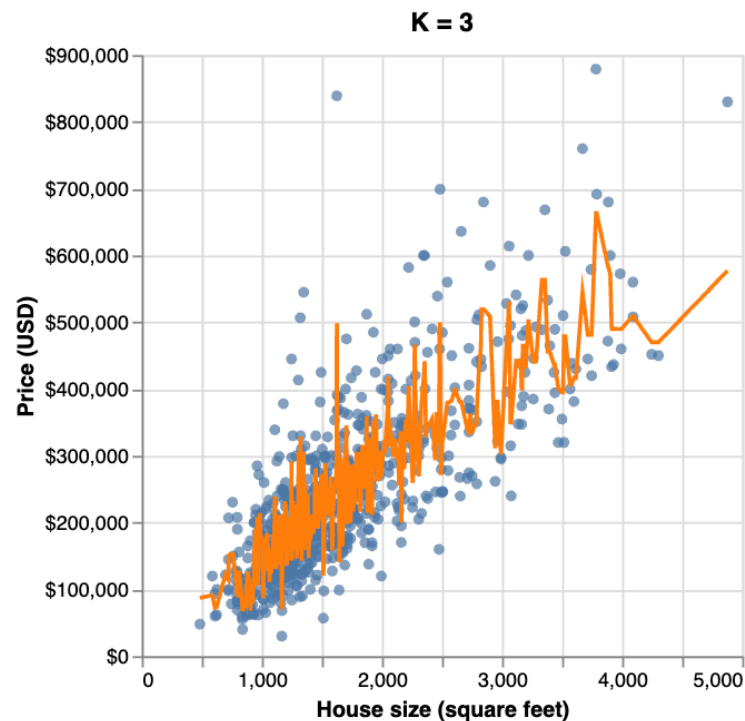
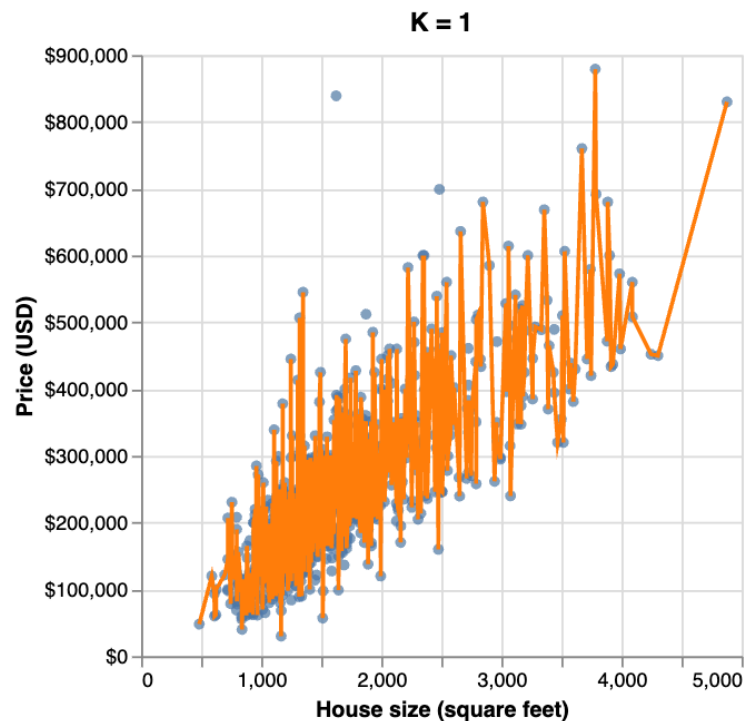
```
1 sacramento_gridsearch.best_params_  
{'kneighborsregressor__n_neighbors': 55}
```

Underfitting and overfitting



The RMSPE values start to get higher after a certain k value

Visualizing different values of k



Evaluating on the test set

RMSPE on the test data

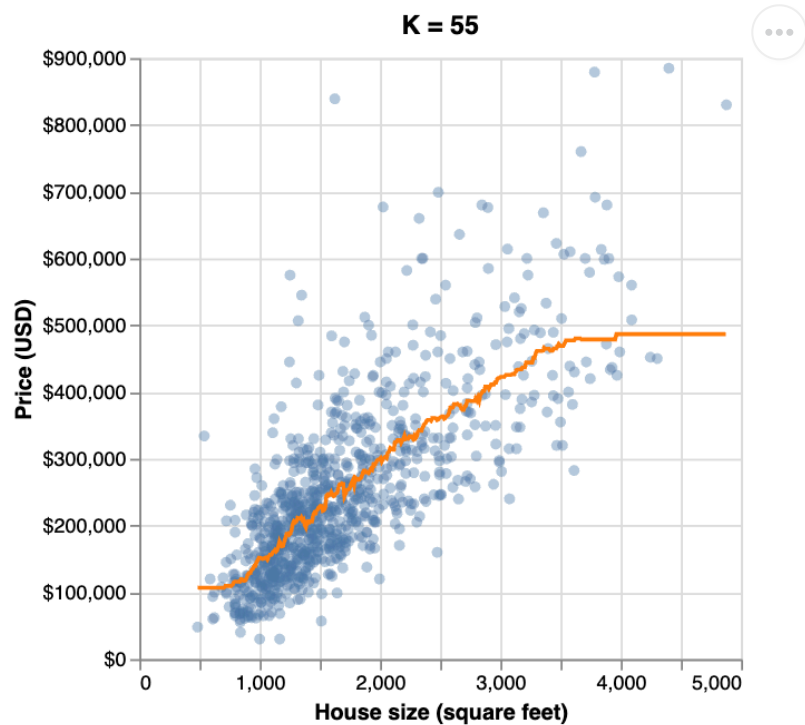
- Retrain the K-NN regression model on the entire training data set using best k

```
1 from sklearn.metrics import mean_squared_error
2
3 sacramento_test['predicted'] = sacramento_gridsearch.predict(sacramento_test)
4 RMSPE = mean_squared_error(
5     y_true=sacramento_test['price'], y_pred=sacramento_test['predicted']
6 ) ** (1 / 2)
7
8 RMSPE
```

```
np.float64(87498.86808211416)
```

Final best K model

Predicted values of house price (orange line) for the final K-NN regression model.



Multivariable K-NN regression

We can use multiple predictors in K-NN regression

Multivariable K-NN regression: Preprocessor

```
1 sacramento_preprocessor = make_column_transformer(  
2     (StandardScaler(), ['sqft', 'beds'])  
3 )  
4 sacramento_pipeline = make_pipeline(sacramento_preprocessor, KNeighborsRegressor())
```

Multivariable K-NN regression: CV

```
1 # create the 5-fold GridSearchCV object
2 param_grid = {
3     'kneighborsregressor__n_neighbors': range(1, 50),
4 }
5
6 sac_r_gridsearch = GridSearchCV(
7     estimator=sacr_pipeline,
8     param_grid=param_grid,
9     cv=5,
10    scoring='neg_root_mean_squared_error',
11 )
12
13 sac_r_gridsearch.fit(
14     sacramento_train[['sqft', 'beds']], sacramento_train['price']
15 )
```

GridSearchCV



best_estimator_: Pipeline

▶ **columntransformer: ColumnTransformer** ?

▶ **standardscaler**

▶ StandardScaler ?

▶ KNeighborsRegressor ?

Multivariable K-NN regression: Best K

```
1 # retrieve the CV scores
2 sacramento = pd.DataFrame(sacr_gridsearch.cv_results_)
3 sacramento['sem_test_score'] = sacramento['std_test_score'] / 5 ** (1 / 2)
4 sacramento['mean_test_score'] = -sacramento['mean_test_score']
5 sacramento = sacramento[
6     [
7         'param_kneighborsregressor__n_neighbors',
8         'mean_test_score',
9         'sem_test_score',
10    ]
11 ].rename(columns={'param_kneighborsregressor__n_neighbors': 'n_neighbors'})
12
13 # show only the row of minimum RMSPE
14 sacramento.nsmallest(1, 'mean_test_score')
```

n_neighbors		mean_test_score	sem_test_score
28	29	85156.027067	3376.143313

Multivariable K-NN regression: Best model

```
1 best_k_sacr_multi = sacr_results["n_neighbors"][sacr_results["mean_test_score"].idxmin()]\n2 min_rmspe_sacr_multi = min(sacr_results["mean_test_score"])
```

Best K

```
1 best_k_sacr_multi
```

```
np.int64(29)
```

Best RMSPE

```
1 min_rmspe_sacr_multi
```

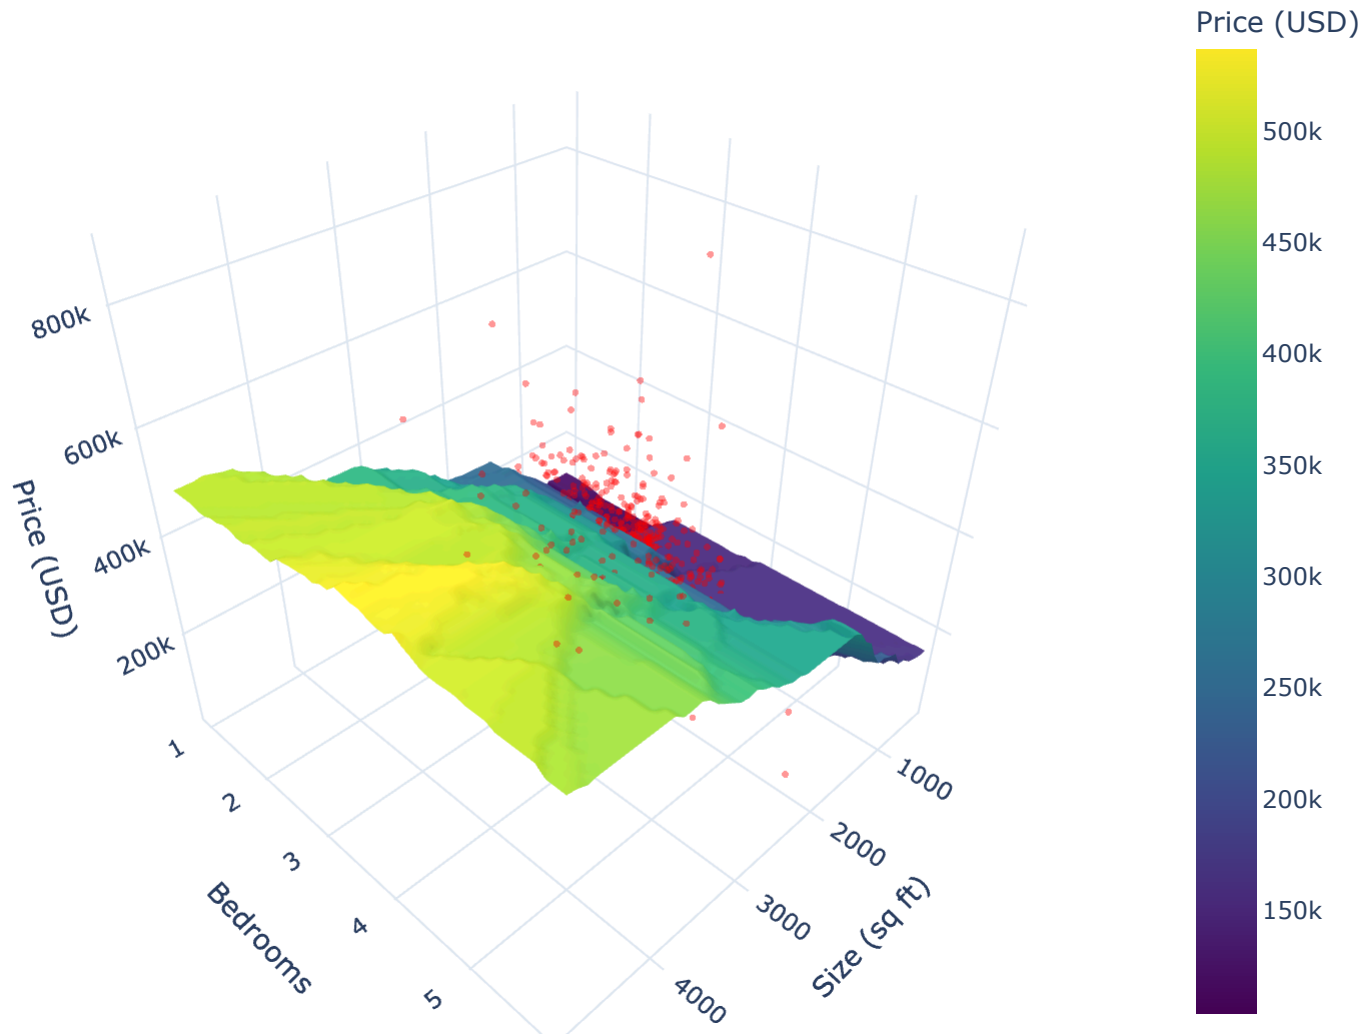
```
85156.02706746716
```

Multivariable K-NN regression: Test data

```
1 sacramento_test["predicted"] = sacr_gridsearch.predict(sacramento_test)
2 RMSPE_mult = mean_squared_error(
3     y_true=sacramento_test["price"],
4     y_pred=sacramento_test["predicted"]
5 )**(1/2)
6
7 RMSPE_mult
```

```
np.float64(85083.2902421959)
```

Multivariable K-NN regression: Visualize



Strengths and limitations of K-NN regression

Strengths:

- simple, intuitive algorithm
- requires few assumptions about what the data must look like
- works well with non-linear relationships (i.e., if the relationship is not a straight line)

Weaknesses:

- very slow as the training data gets larger
- may not perform well with a large number of predictors
- may not predict well beyond the range of values input in your training data

Linear Regression

- Addresses the limitations from KNN regression
- provides an interpretable mathematical equation that describes the relationship between the predictor and response variables
- Create a straight line of best fit through the training data

Note

Logistic regression is the linear model we can use for binary classification

Sacramento real estate

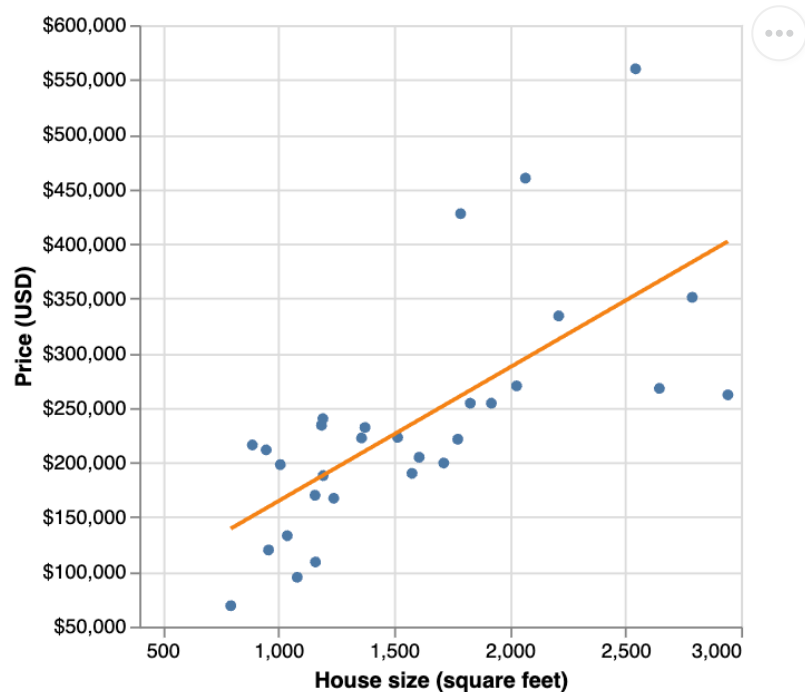
```
1 import pandas as pd
2
3 sacramento = pd.read_csv('data/sacramento.csv')
4
5 np.random.seed(42)
6 small_sacramento = sacramento.sample(n=30)
7
8 print(small_sacramento)
```

	city	zip	beds	baths	sqft	type	price	latitude	\
829	SACRAMENTO	z95824	3	1.0	1161	Residential	109000	38.511893	
70	ELK_GROVE	z95624	4	2.0	1715	Residential	199500	38.440760	
..	
909	SACRAMENTO	z95828	3	1.0	888	Residential	216021	38.508217	
265	SACRAMENTO	z95835	4	3.0	2030	Residential	270000	38.671807	

	longitude
829	-121.457676
70	-121.385792
..	...
909	-121.411207
265	-121.498274

[30 rows x 9 columns]

Sacramento real estate: best fit line



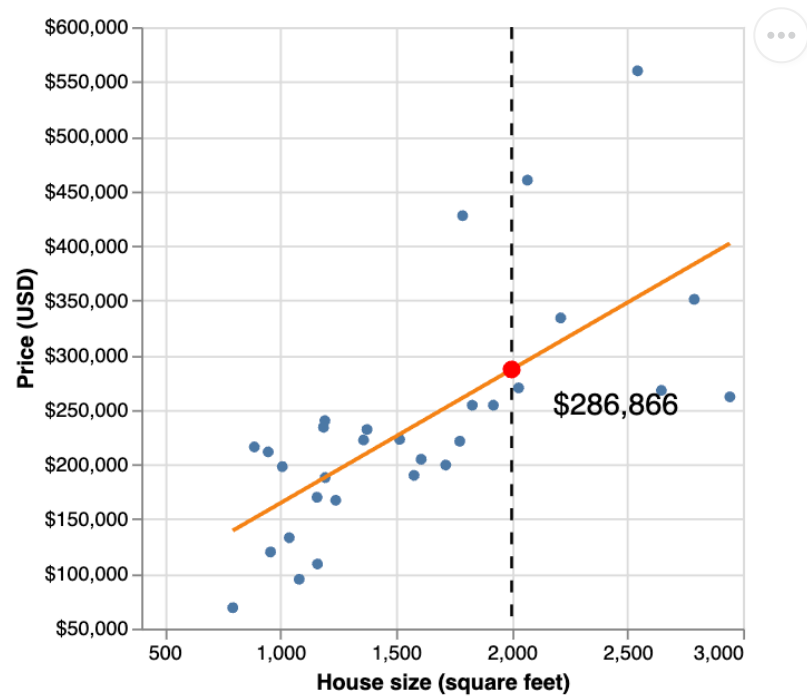
The equation for the straight line is:

$$\text{house sale price} = \beta_0 + \beta_1 \cdot (\text{house size}),$$

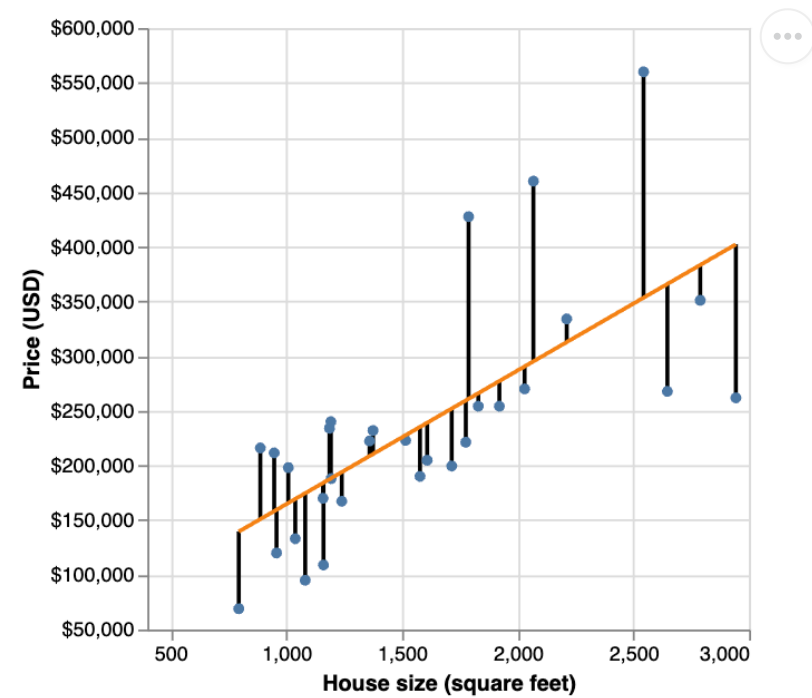
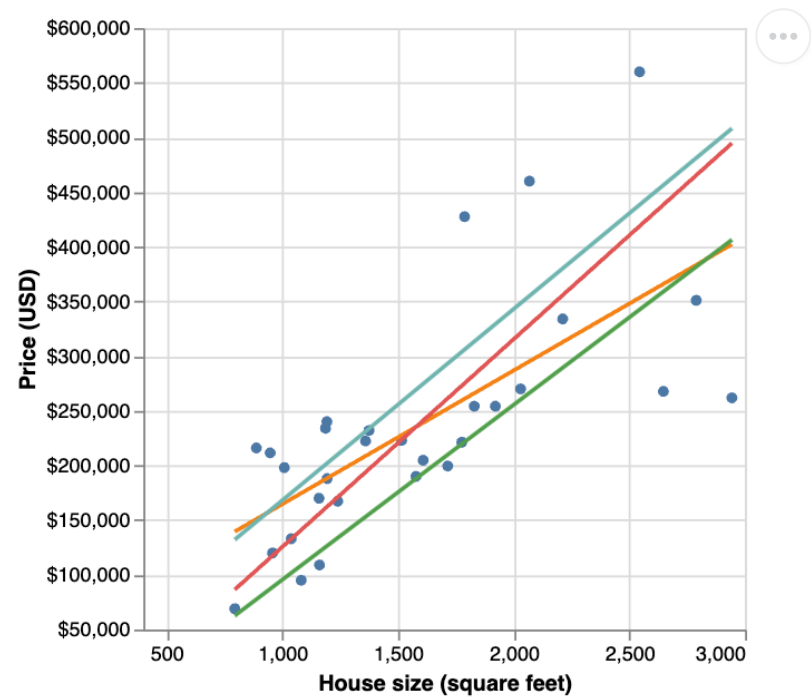
where

- β_0 is the *vertical intercept* of the line (the price when house size is 0)
- β_1 is the *slope* of the line (how quickly the price increases as you increase house size)

Sacramento real estate: Prediction



What makes the best line?



Linear regression in Python

The `scikit-learn` pattern still applies:

- Create a training and test set
- Instantiate a model
- Fit the model on training
- Use model on testing set

Linear regression: Train test split

```
1 import numpy as np
2 import altair as alt
3 import pandas as pd
4 from sklearn.model_selection import train_test_split
5 from sklearn.linear_model import LinearRegression
6 from sklearn.metrics import mean_squared_error
7 from sklearn import set_config
8
9 # Output dataframes instead of arrays
10 set_config(transform_output='pandas')
11
12 np.random.seed(1)
13
14 sacramento = pd.read_csv('data/sacramento.csv')
15
16 sacramento_train, sacramento_test = train_test_split(
17     sacramento, train_size=0.75
18 )
```

Linear regression: Fit the model

```
1 # fit the linear regression model
2 lm = LinearRegression()
3 lm.fit(
4     sacramento_train[['sqft']], # A single-column data frame
5     sacramento_train['price'], # A series
6 )
7
8 # make a dataframe containing slope and intercept coefficients
9 results_df = pd.DataFrame({'slope': [lm.coef_[0]], 'intercept': [lm.intercept_]})
10 print(results_df)
```

	slope	intercept
0	137.285652	15642.309105

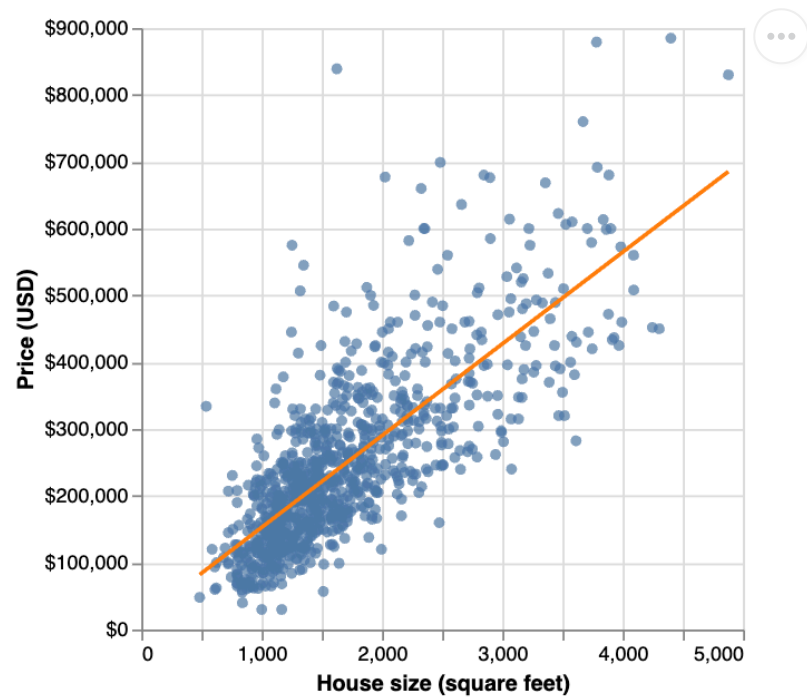
house sale price = $137.29 + 15642.31 \cdot (\text{house size})$.

Linear regression: Predictions

```
1 # make predictions
2 sacramento_test["predicted"] = lm.predict(sacramento_test[["sqft"]])
3
4 # calculate RMSPE
5 RMSPE = mean_squared_error(
6     y_true=sacramento_test["price"],
7     y_pred=sacramento_test["predicted"]
8 )**(1/2)
9
10 RMSPE
```

```
np.float64(85376.59691629931)
```

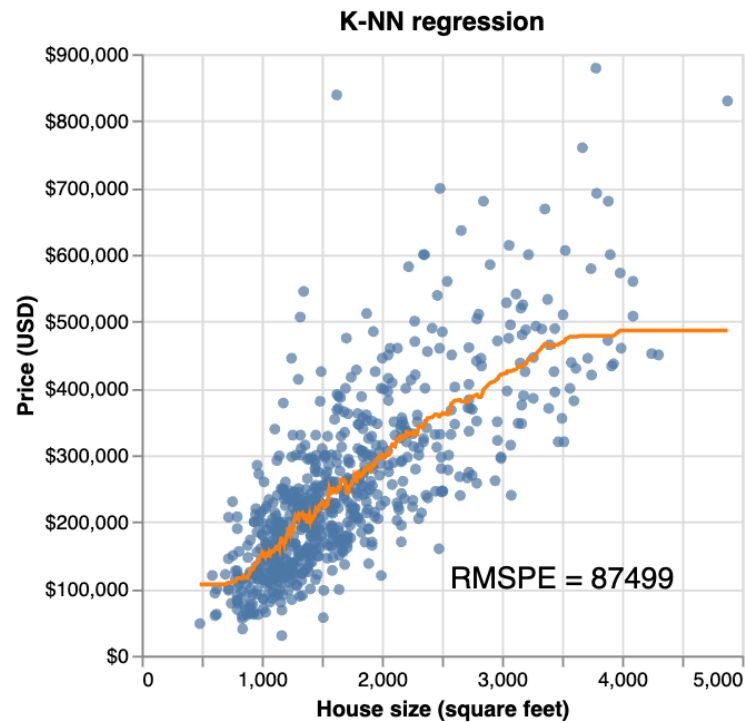
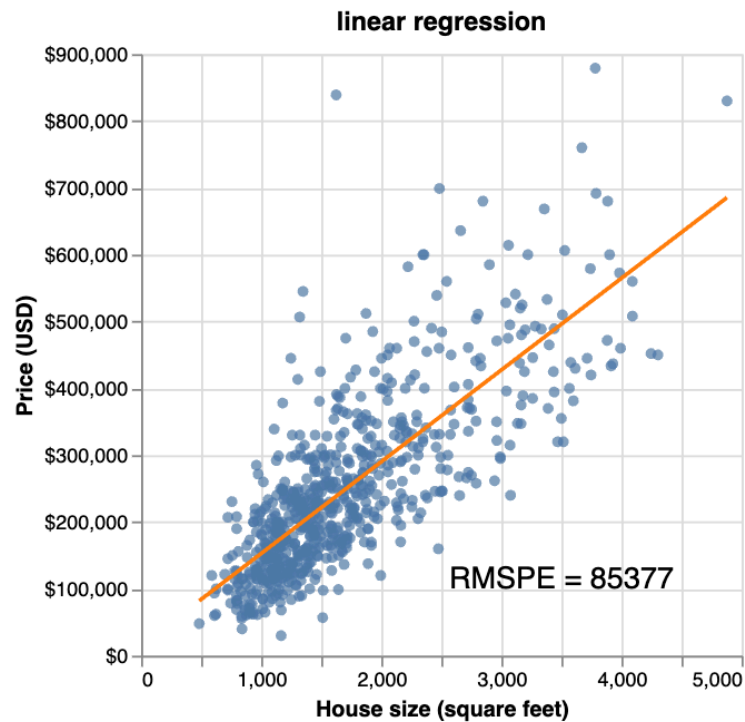
Linear regression: Plot



Standardization

- We did not need to standarize like we did for KNN.
- In KNN standarization is mandatory,
- In linear regression, if we standarize, we convert all the units to unit-less standard deviations
- Standarization in linear regression does not change the fit of the model
 - It will change the coefficients

Comparing simple linear and K-NN regression



Multivariable linear regression

More predictor variables!

- More does not always mean better
- We will not cover variable selection in this workshop
- Will talk about categorical predictors later in the workshop

Sacramento real estate: 2 predictors

```
1 mlm = LinearRegression()
2 mlm.fit(sacramento_train[['sqft', 'beds']], sacramento_train['price'])
3
4 sacramento_test['predicted'] = mlm.predict(sacramento_test[['sqft', 'beds']])
```

Sacramento real estate: Coefficients

Coefficients

```
1 mlm.coef_
```

```
array([ 154.59235377, -20333.43213798])
```

Intercept

```
1 mlm.intercept_
```

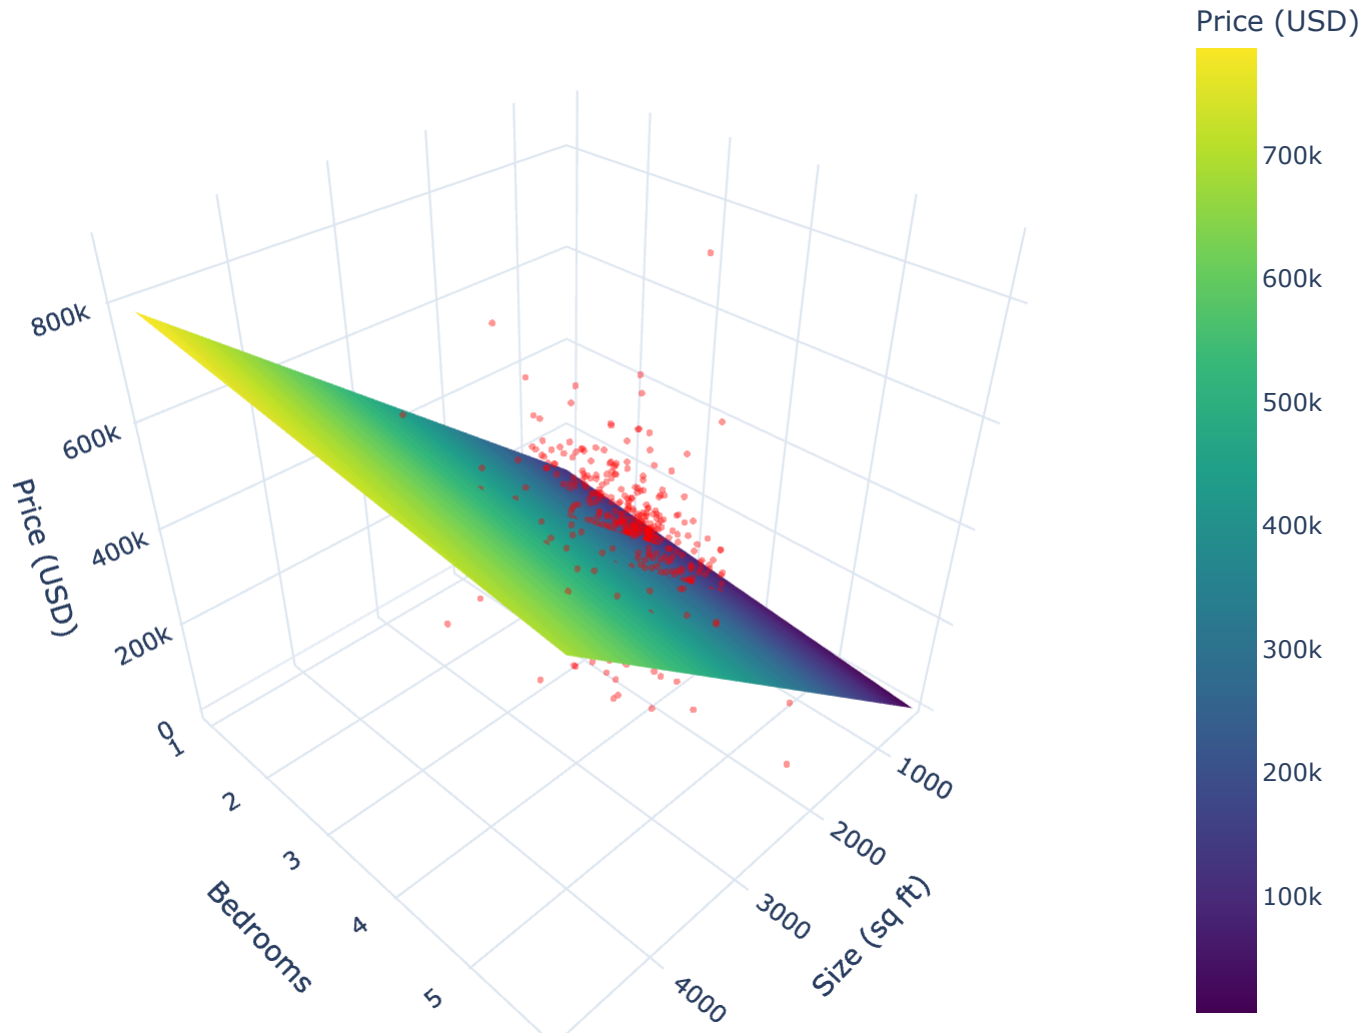
```
np.float64(53180.26906624224)
```

$$\text{house sale price} = \beta_0 + \beta_1 \cdot (\text{house size}) + \beta_2 \cdot (\text{number of bedrooms}),$$

where:

- β_0 is the *vertical intercept* of the hyperplane (the price when both house size and number of bedrooms are 0)
- β_1 is the *slope* for the first predictor (how quickly the price increases as you increase house size)
- β_2 is the *slope* for the second predictor (how quickly the price increases as you increase the number of bedrooms)

More variables make it harder to visualize



Sacramento real estate: mlm rmspe

```
1 lm_mult_test_RMSPE = mean_squared_error(  
2     y_true=sacramento_test['price'], y_pred=sacramento_test['predicted']  
3 ) ** (1 / 2)  
4  
5 lm_mult_test_RMSPE
```

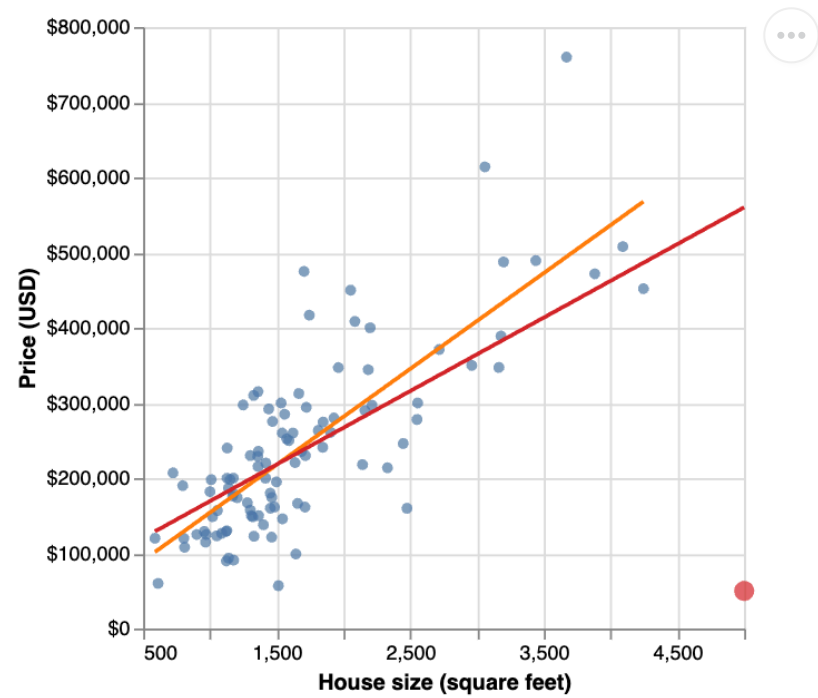
```
np.float64(82331.04630202598)
```

Outliers and Multicollinearity

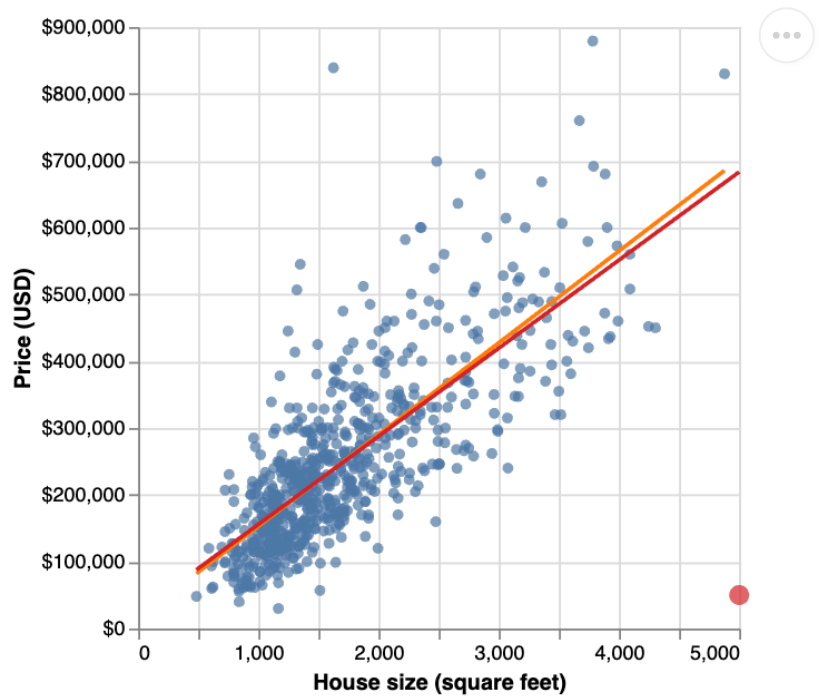
- Outliers: extreme values that can move the best fit line
- Multicollinearity: variables that are highly correlated to one another

Outliers

Subset

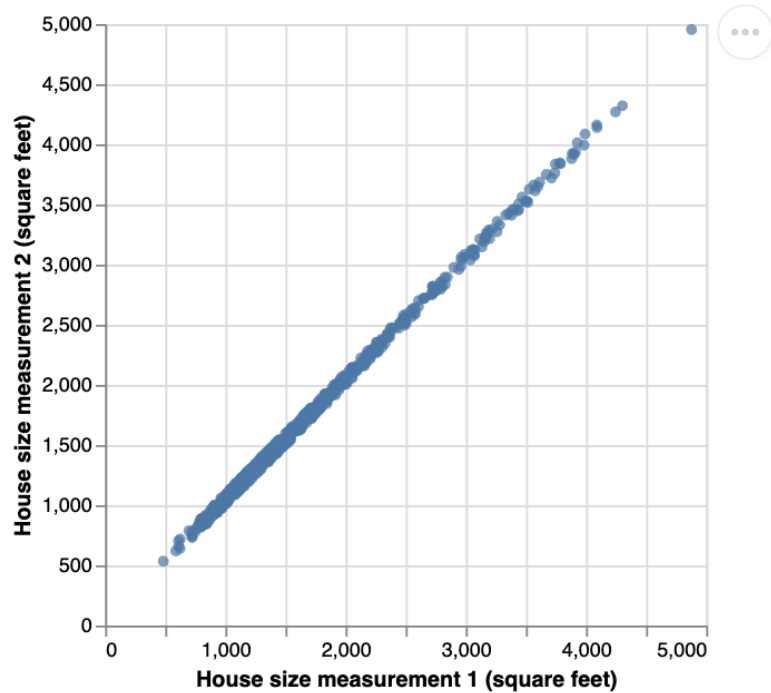


Full data



Multicollinearity

plane of best fit has regression coefficients that are very sensitive to the exact values in the data



- The [Regression I: K-nearest neighbors](#) and [Regression II: linear regression](#) chapters of Data Science: A First Introduction (Python Edition) by Tiffany Timbers, Trevor Campbell, Melissa Lee, Joel Ostblom, Lindsey Heagy contains all the content presented here with a detailed narrative.
- The [scikit-learn website](#) is an excellent reference for more details on, and advanced usage of, the functions and packages in this lesson. Aside from that, it also offers many useful [tutorials](#) to get you started.
- *[An Introduction to Statistical Learning](#)* by Gareth James Daniela Witten Trevor Hastie, and Robert Tibshirani provides a great next step in the process of learning about classification. Chapter 3 discusses linear regression in more depth. As well as how it compares to K-nearest neighbors.

References

Thomas Cover and Peter Hart. Nearest neighbor pattern classification. IEEE Transactions on Information Theory, 13(1):21–27, 1967.

Evelyn Fix and Joseph Hodges. Discriminatory analysis. nonparametric discrimination: consistency properties. Technical Report, USAF School of Aviation Medicine, Randolph Field, Texas, 1951.